

Parsers in the Compilation Process for RISC-V Architectures

Table of Contents

1. Introduction
 2. The Connection Between Scanners and Parsers
 3. Expressing Syntax: Grammars and Parsing
 4. Top-Down Parsing
 5. Bottom-Up Parsing
 6. Practical Issues in Parsing
 - Error Recovery
 - Handling Unary Operations
 - Context-Sensitive Ambiguity
 7. Parsers in LLVM and RISC-V Compilation
 8. Conclusion
 9. References
-

1. Introduction

Parsing is a fundamental step in the compilation process. It takes the tokenized input from the scanner and structures it into a hierarchical form—usually a **parse tree** or an **abstract syntax tree (AST)**—which represents the syntactic structure of the source program. This process ensures that the program adheres to the grammatical rules of the programming language and prepares it for semantic analysis and code generation.

For RISC-V architectures, parsing is particularly important in generating efficient low-level code. The way constructs are parsed and transformed impacts the generated assembly instructions, influencing performance and optimization opportunities within an LLVM-based compiler toolchain. The efficiency of parsing can determine how well high-level constructs translate into optimized machine code, ensuring minimal execution overhead and effective utilization of the RISC-V instruction set.

Parsing also plays a role in software security and reliability. A well-designed parser ensures that syntactic errors, ambiguities, and unintended behavior are detected early, preventing possible vulnerabilities. Many modern compilers integrate defensive parsing techniques to prevent attacks such as injection-based exploits in interpreters and compilers.

Furthermore, parsers facilitate multiple stages of the compilation process, including macro expansion, type checking, and intermediate representation (IR) generation. The ability to recognize and structure program constructs influences the effectiveness of compiler optimizations such as **constant folding**, **dead code elimination**, and **loop unrolling**—all of which are crucial for improving execution efficiency on RISC-V processors.

Historically, parsing techniques have evolved significantly, from simple hand-written recursive descent parsers to highly efficient, table-driven LR parsers. These advances have enabled modern compilers to handle increasingly complex programming languages with rich syntactic constructs. In the context of RISC-V, parsing efficiency is especially crucial due to the architecture's focus on **reduced instruction set computing (RISC)** principles, which emphasize streamlined and predictable execution paths.

Moreover, parsing affects not just compilation but also software tooling, including **static analyzers, interpreters, formatters, and linting tools**. A robust parsing strategy ensures that such tools provide accurate feedback, improving the software development workflow for RISC-V applications.

Given its pivotal role in compiler technology, parsing requires a deep understanding of grammar structures, parsing algorithms, and error recovery mechanisms, all of which will be explored in this document.

2. The Connection Between Scanners and Parsers

A **scanner (lexical analyzer)** processes the raw source code and converts it into a stream of **tokens**. These tokens are then fed into the **parser (syntax analyzer)**, which organizes them according to grammatical rules.

While scanners perform simple lexical analysis by recognizing keywords, identifiers, literals, and operators, parsers go further by establishing hierarchical relationships between these elements. The integration between scanners and parsers must be efficient to prevent performance bottlenecks. In many compiler toolchains, scanner performance is improved by techniques like lookahead buffering, finite automata optimizations, and parallel tokenization strategies.

One critical aspect of the scanner-parser relationship is handling whitespace, comments, and preprocessing directives. While these elements are typically discarded by the scanner, they may influence parsing decisions, especially in languages with complex macro systems such as C. For instance, the C preprocessor modifies the source code before tokenization, introducing challenges in preserving accurate error locations and source mapping in the parsing stage.

Another important consideration is the role of **lookahead** in parsing. Many parsers rely on a limited number of lookahead tokens to determine which production rule to apply next. The scanner must be capable of buffering tokens ahead of time, ensuring that the parser can resolve ambiguous constructs efficiently. This is particularly important in RISC-V compilers when parsing assembly-like constructs embedded within C or C++ code, where specific token sequences require precise parsing strategies.

Additionally, modern compilers, including those targeting RISC-V, often implement **lazy lexical analysis**, where tokenization occurs on demand rather than upfront for the entire file. This approach reduces memory overhead and allows parsers to adjust dynamically based on language context (e.g., distinguishing between types and variable names in C++ template

parsing). By integrating scanning and parsing closely, compiler designers can optimize performance and maintain parsing flexibility for evolving language features.

Example: Tokenizing and Parsing a C Expression


Consider the following simple C expression:

```
x = a + b * c;
```

A scanner might produce the following token sequence:

```
IDENTIFIER(x) ASSIGN IDENTIFIER(a) PLUS IDENTIFIER(b) MULT IDENTIFIER(c)
SEMICOLON
```

The parser then processes these tokens to build an AST:



This AST is later translated into LLVM Intermediate Representation (LLVM IR) and optimized for the RISC-V backend. This structured representation enables further compiler optimizations, such as instruction reordering, register allocation, and loop unrolling, which are crucial for high-performance execution on RISC-V processors.

3. Expressing Syntax: Grammars and Parsing

A parser is guided by a formal **context-free grammar (CFG)**, often expressed in **Backus-Naur Form (BNF)**. A simple grammar for arithmetic expressions in C could be:

```
E → E + T | E - T | T
T → T * F | T / F | F
F → (E) | id | num
```

This grammar describes valid sequences of operations and operands. However, it has left recursion, which must be removed for certain parsing techniques like LL(1).

Grammar design affects how easily a parser can process a language. Ambiguities, left recursion, and unnecessary complexities can make parsing inefficient. Optimizing a grammar to fit a parsing method improves compilation speed and accuracy. Furthermore, some grammars require additional transformation steps, such as removing common left factors or factoring out recursion, to fit specific parsing strategies.

Beyond basic CFGs, real-world programming languages often require **attribute grammars** that include semantic rules alongside syntax rules. These attributes help resolve

issues such as type checking, variable scope determination, and function signature validation during parsing.


Parsing strategies also rely on **operator precedence** to disambiguate expressions. Without explicit precedence rules, expressions such as `a + b * c` could be parsed incorrectly. Many parsers use **precedence climbing** or **operator-precedence parsing** techniques to enforce standard arithmetic evaluation rules, ensuring that multiplication and division are evaluated before addition and subtraction.

Example: Handling Operator Precedence in a Parser


A typical operator-precedence table for C expressions might look like this:

Operator	Associativity	Precedence Level
<code>*, /</code>	Left	2
<code>+, -</code>	Left	1

Using this table, a parser ensures that `a + b * c` is parsed as:



rather than:



For compilers targeting RISC-V, grammar definition plays a critical role in efficiently translating high-level operations into minimal and well-optimized machine instructions. Advanced parsing strategies, such as **packrat parsing** or **GLR (Generalized LR) parsing**, can be employed when dealing with complex language constructs, particularly in Just-In-Time (JIT) compilers that optimize code on the fly.

4. Top-Down Parsing

Top-down parsers start at the root of the parse tree and work downwards. These parsers are particularly well-suited for smaller, well-structured languages but struggle with more complex language constructs found in modern programming languages.

Recursive Descent Parsing

This method uses recursive functions to parse an input stream. Consider a recursive descent parser for arithmetic expressions:

```
void E() {
    T();
    while (lookahead == '+' || lookahead == '-') {
        char op = lookahead;
        match(op);
        T();
    }
}

void T() {
    F();
    while (lookahead == '*' || lookahead == '/') {
        char op = lookahead;
        match(op);
        F();
    }
}
```

This approach is simple but struggles with left-recursive grammars, requiring transformation before parsing.

Recursive descent parsers are commonly used in hand-written compilers because they are straightforward to implement and debug. However, they require significant manual effort to handle complex syntax and ambiguities, making them less scalable for large languages.

Eliminating Left Recursion

One major limitation of recursive descent parsing is that it cannot handle **left-recursive grammars**, which cause infinite recursion. A left-recursive rule like:

$$E \rightarrow E + T \mid T$$

needs to be rewritten in a **right-recursive** or **iterative** form:

$$\begin{aligned} E &\rightarrow T E' \\ E' &\rightarrow + T E' \mid \varepsilon \end{aligned}$$

This transformation allows a recursive descent parser to process input without infinite recursion.

Predictive Parsing and LL(1) Parsers

Predictive parsing eliminates the need for backtracking by using a parsing table and a lookahead token to decide the next step. LL(1) parsers are a common form of predictive parsers that use one-token lookahead. The parsing table for a given grammar can be constructed manually or using tools such as `antlr`.

A typical LL(1) parsing table might look like this for a simple grammar:

Non-terminal	id	+	*	()	\$
E	$E \rightarrow TE'$			$E \rightarrow TE'$		
E'		$E' \rightarrow +TE'$			$E' \rightarrow \varepsilon$	$E' \rightarrow \varepsilon$
T	$T \rightarrow FT'$			$T \rightarrow FT'$		
T'		$T' \rightarrow \varepsilon$	$T' \rightarrow *FT'$		$T' \rightarrow \varepsilon$	$T' \rightarrow \varepsilon$
F	$F \rightarrow \text{id}$			$F \rightarrow (E)$		

LL(1) parsing is useful for simple grammars but struggles with left-recursion and ambiguous rules, making it unsuitable for many real-world languages without preprocessing.

Handling Ambiguities in Top-Down Parsing

Many programming languages introduce ambiguities that must be resolved during parsing. A common example is the **dangling else problem** in `if-else` statements:

```
if (x > 0)
    if (y > 0)
        printf("Both positive");
    else
        printf("Only x is positive");
```

An LL(1) parser struggles with this because it cannot decide whether the `else` belongs to the inner or outer `if`. Solutions include:

- Using **explicit grammar rules** to enforce correct parsing.
- Leveraging **indentation-based syntax**, as seen in Python.
- Using **parser precedence** to associate `else` with the nearest unmatched `if`.

Example: Parsing a Simple Expression

Consider the input:

`x + y * z`

Using a top-down parser, this expression is parsed as:



This structure ensures correct precedence handling, essential for correct evaluation in compiled RISC-V code.

Advantages and Disadvantages of Top-Down Parsing

Advantages:

1. **Easy to Implement:** Recursive descent parsers are intuitive and simple to write.
2. **Good for Small Languages:** Well-suited for DSLs and configuration file parsing.
3. **Predictability:** LL(1) parsers have predictable behavior and no shift-reduce conflicts.

Disadvantages:

1. **Cannot Handle Left Recursion:** Requires grammar transformation before parsing.
2. **Limited in Power:** LL parsers cannot handle a wide range of real-world grammars without extra lookahead or transformations.
3. **Backtracking Overhead:** Non-LL(1) grammars require backtracking, which is inefficient.

Top-Down Parsing in RISC-V Compilation

While bottom-up parsing is more common in industrial-strength compilers, top-down parsing is sometimes used in **domain-specific languages (DSLs)**, **simple scripting languages**, and **interpreters** for embedded systems. In RISC-V toolchains, top-down parsing may be utilized in small utility compilers or specialized code transformations where performance and complexity are less critical.

5. Bottom-Up Parsing

Bottom-up parsers construct the parse tree from the leaves up. **LR parsers**, such as those generated by `yacc` or `bison`, are widely used in compiler construction.

Shift-Reduce Parsing

One of the fundamental techniques in bottom-up parsing is **shift-reduce parsing**, where a stack is used to hold symbols, and reductions are applied based on the grammar rules. Unlike top-down parsing, bottom-up parsing does not require left-factored grammars, making it more suitable for handling complex language constructs.

Shift-reduce parsing involves two key actions:

- **Shift:** Push the next input token onto the stack.
- **Reduce:** Replace a sequence of tokens in the stack with a corresponding non-terminal based on a grammar rule.

*Example: Parsing $a + b * c$ with Shift-Reduce Parsing*

For the input $a + b * c$, the parsing actions would be:

Stack	Input	Action
a	+ b * c	Shift
a +	b * c	Shift
a + b	* c	Reduce ($T \rightarrow b$)
a + T	* c	Shift
a + T *	c	Shift
a + T * c		Reduce ($T \rightarrow T * F$)
a + T		Reduce ($E \rightarrow E + T$)

The final reduction step yields a fully parsed expression, ready for LLVM IR generation.

LR Parsing and Its Variants

LR parsers are widely used for their power and ability to handle complex grammars. Several variants exist:

- **SLR (Simple LR):** The simplest LR parsing technique, using follow sets for reductions.
- **LR(1) (Canonical LR):** Uses lookahead tokens to resolve ambiguities and provides greater parsing power.
- **LALR (Look-Ahead LR):** A more memory-efficient variation of LR(1), used in tools like `bison`.

LR parsing follows a table-driven approach, where the parser maintains a **parsing table** that defines the possible actions based on the current state and input token.

Example: LR Parsing Table for a Simple Grammar

Consider the grammar:

$E \rightarrow E + T \mid T$
 $T \rightarrow T * F \mid F$
 $F \rightarrow (E) \mid id$

A simplified LR parsing table for this grammar might look like:

State	id	+	*	()	\$	Action
0	S3			S4			
1		S5				Accept	
2		R2	S6		R2	R2	
3		R4	R4		R4	R4	
4	S3			S4			
5	S3			S4			
6	S3			S4			

Each **S** (shift) and **R** (reduce) entry specifies an action, helping guide the parser through the input sequence.

Advantages of Bottom-Up Parsing

1. **Handles a Larger Class of Grammars:** Unlike LL parsers, LR parsers can handle left-recursive grammars, which are common in real-world programming languages.
2. **More Powerful Parsing Decisions:** By postponing reductions, LR parsers make parsing decisions with more context, avoiding incorrect derivations.
3. **Efficient Table-Driven Parsing:** With precomputed parsing tables, bottom-up parsers can be implemented efficiently, making them suitable for industrial compilers.

Bottom-Up Parsing in RISC-V Compilation

For RISC-V compilation, bottom-up parsing ensures that complex expressions and language constructs are efficiently translated into an optimal sequence of instructions. This method plays a key role in converting **higher-level constructs into efficient machine-level operations** that match the constraints of the RISC-V instruction set.

Moreover, efficient parsing contributes to improved **compiler optimizations**, such as constant folding, dead code elimination, and instruction selection. Since RISC-V follows a **reduced instruction set computing (RISC)** approach, ensuring efficient translation of parsed constructs into minimal instructions is crucial for achieving high performance.

Bottom-up parsing is widely used in modern compiler toolchains, including **LLVM**, where production-grade parsers employ **LALR(1)** and **GLR (Generalized LR)** parsing strategies for handling complex language features while maintaining efficiency.

6. Practical Issues in Parsing

Parsing in real-world compilers involves several practical challenges that must be addressed to ensure robustness, efficiency, and correctness. While theoretical parsing methods provide a structured way to process syntax, real-world constraints introduce complexities that compilers must handle effectively. Three key issues in practical parsing include **error recovery**, **handling unary operations**, and **context-sensitive ambiguities**.

Error Recovery

Errors in parsing can arise from syntax mistakes, missing tokens, or incorrect token sequences. A well-designed compiler should handle these errors gracefully, providing useful feedback to the programmer.

Techniques for Error Recovery

1. **Panic Mode Recovery:** When an error is encountered, the parser discards tokens until it finds a synchronization point, such as a semicolon (;) or closing brace (}).

2. **Phrase-Level Recovery:** The parser attempts to correct the erroneous phrase by inserting or deleting tokens, making minor modifications to continue parsing.
3. **Error Productions:** The grammar is extended to include common error patterns, allowing specific recovery actions.
4. **Global Correction:** The parser finds the minimal set of modifications needed to correct the program, although this approach is computationally expensive.

For RISC-V compilers, error recovery mechanisms ensure that syntax errors do not propagate into later compilation stages, preventing faulty assembly generation.

Handling Unary Operations

Unary operators, such as negation ($-x$) or logical NOT ($!x$), introduce parsing challenges since $-$ and $!$ can also be binary operators in some contexts.

Example Ambiguity

```
int x = -5;
int y = a - -b;
```

In the second line, $a - -b$ involves two consecutive $-$ operators, which must be correctly interpreted as subtraction and negation.

Solutions

- Modify the grammar to distinguish between **prefix unary operators** and **binary operators**.
- Use **operator precedence rules** to disambiguate expressions.
- Implement **symbol table lookups** to determine context-sensitive meaning.

Correct handling of unary operations ensures that mathematical expressions are parsed correctly and efficiently translated into RISC-V assembly instructions.

Context-Sensitive Ambiguity

Most parsing techniques assume **context-free grammars**, but real programming languages often require **context-sensitive analysis**.

Example: Type vs. Variable Ambiguity

```
int a;
int b(int);
```

The parser must determine whether $b(int)$ is a function declaration or a multiplication operation with parentheses.

Resolution Strategies

- **Two-pass parsing:** A preliminary pass identifies declarations before parsing expressions.
- **Symbol table integration:** Using semantic analysis during parsing helps resolve ambiguous constructs.

- **Lookahead tokens:** Checking multiple tokens ahead allows early disambiguation.

In the context of RISC-V, handling context-sensitive constructs correctly ensures that function calls, memory access patterns, and optimizations align with the architecture's constraints.

By addressing these practical issues, modern parsers provide robust and efficient compilation, bridging the gap between high-level programming constructs and optimized RISC-V machine code.

7. Parsers in LLVM and RISC-V Compilation

LLVM (Low-Level Virtual Machine) is a widely used compiler framework that provides a modular approach to building language frontends, optimizers, and backend code generators. The LLVM infrastructure includes robust parsing mechanisms that play a key role in translating high-level source code into optimized RISC-V assembly.

Parsing in the LLVM Compiler Frontend

LLVM-based compilers, such as **Clang**, rely on a structured parsing pipeline:

1. **Lexical Analysis:** The scanner tokenizes the source code into recognizable symbols.
2. **Parsing (Syntax Analysis):** The parser builds an **Abstract Syntax Tree (AST)** based on the grammar of the language.
3. **Semantic Analysis:** The AST undergoes type checking, name resolution, and validation.
4. **Intermediate Representation (IR) Generation:** The validated AST is converted into LLVM IR for optimization and backend processing.

LLVM primarily employs **recursive descent parsing** for C and C++ but integrates **table-driven parsing techniques** for more complex language features. The ability to analyze code structurally allows LLVM to generate efficient, optimized IR that can be transformed into RISC-V machine code.

Example: LLVM IR Generation from C Code

Consider the following C function:

```
int add(int a, int b) {  
    return a + b;  
}
```

The Clang frontend parses this into an AST, which is then lowered into LLVM IR:

```
define i32 @add(i32 %a, i32 %b) {
entry:
    %sum = add i32 %a, %b
    ret i32 %sum
}
```

This IR is then optimized and passed to the RISC-V backend.

Parsing Challenges in RISC-V Compilation

1. Instruction Selection and Parsing Constraints

Since RISC-V is a **reduced instruction set computing (RISC)** architecture, parsing and code generation must ensure that complex operations are broken into simple instructions. Unlike x86, which has variable-length instructions, RISC-V requires strict adherence to fixed 32-bit or compressed 16-bit instruction formats. This constraint impacts parsing decisions when generating assembly from IR.

2. Handling Inline Assembly and Pseudo-Instructions

Many compilers support **inline assembly**, allowing programmers to write RISC-V instructions directly within C or C++. Parsing such constructs correctly requires the parser to distinguish between standard syntax and inline assembly blocks.

Example:

```
asm("add t0, t1, t2");
```

This requires special handling in LLVM's frontend to ensure correct instruction selection and register allocation.

3. Optimized Parsing for Vector and Floating-Point Instructions

RISC-V extensions introduce additional complexity in parsing. For example, **RISC-V Vector Instructions** allow SIMD operations, which must be parsed with awareness of vector registers and element widths.

```
vadd.vv v1, v2, v3 # Vector addition
```

LLVM's parsing framework must handle these constructs while ensuring compliance with RISC-V's modular extension model.

LLVM's Role in RISC-V Backend Code Generation

Once parsing and optimization are complete, LLVM generates **RISC-V assembly**. For instance, the earlier `add` function may be compiled to:

```
add:
    add a0, a0, a1
    ret
```

Parsing and IR generation play a critical role in ensuring that high-level language constructs are efficiently translated into these low-level RISC-V instructions, optimizing for performance and energy efficiency.

8. Conclusion

Parsing is a fundamental component of compiler design, bridging the gap between raw source code and its structured representation in an Abstract Syntax Tree (AST) or Intermediate Representation (IR). In the context of **RISC-V compilation**, efficient parsing ensures that high-level constructs are correctly interpreted and transformed into optimized machine instructions.

Top-down and bottom-up parsing methods provide distinct advantages, with **LL parsers** being easier to implement and **LR parsers** offering greater parsing power. Practical issues such as **error recovery**, **unary operations**, and **context-sensitive ambiguities** play a crucial role in real-world compiler implementations.

LLVM serves as a robust framework for RISC-V compilers, leveraging structured parsing techniques to generate efficient IR and ultimately optimized assembly code. By addressing **instruction constraints**, **inline assembly parsing**, and **vector instruction handling**, LLVM enables smooth translation from high-level languages to the RISC-V architecture.

Understanding parsing techniques is not only crucial for compiler developers but also for engineers working on **domain-specific languages (DSLs)**, **static analysis tools**, and **performance-critical applications** that rely on precise code generation.

Future advancements in parsing technology, such as **machine-learning-driven grammar inference** and **adaptive parsing strategies**, may further improve compiler efficiency and flexibility, making parsing an ever-evolving field in computing.

9. References

- Cooper, K., & Torczon, L. (2011). *Engineering a Compiler*. Elsevier.
- Aho, A. V., Lam, M. S., Sethi, R., & Ullman, J. D. (2006). *Compilers: Principles, Techniques, and Tools* (2nd ed.). Pearson.
- LLVM Documentation: <https://llvm.org/docs/>
- The RISC-V Instruction Set Manual: <https://riscv.org/technical/specifications/>
- Fraser, C. W., & Hanson, D. R. (1995). *A Retargetable C Compiler: Design and Implementation*. Benjamin-Cummings.
- Grune, D., Bal, H., Jacobs, C., & Langendoen, K. (2012). *Modern Compiler Design*. Springer.